

Scale sanity check (pre-norm), per-model normalised: factual leads $\sigma_{\max}$ ONLY at 0.5B; coding leads at 1.5B/3B

